

# Diffusion Models: Theory and Applications

## Lecture 5: Sampling from Trained Diffusion Models - From Noise to Data


Shubham Chatterjee

Missouri University of Science and Technology, Department of Computer Science

June 11, 2025

We mastered the ELBO framework and noise prediction...

**But how do we actually USE these trained models?** 

- ✓ **ELBO decomposition:** Three interpretable forces
- ✓ **Noise prediction:** Simple MSE training objective
- ✓ **Training algorithm:** From complex theory to practical implementation
-  **Today:** Turn trained models into sample generators!

## 🎯 How do we go from pure noise to realistic data?

### The Challenge:

- We have a trained noise predictor  $\epsilon_{\theta}(\mathbf{x}_t, t)$
- Need to reverse the destruction process
- Want high quality AND reasonable speed

### Two Main Approaches:

- **DDPM**: Stochastic, high quality, slow
- **DDIM**: Deterministic, fast, still high quality

**Sampling = Where theory meets practice!**

# The Fundamental Sampling Challenge ?

**We've trained our network to predict noise...**

What we have 

- Trained network:  $\epsilon_{\theta}(\mathbf{x}_t, t)$  predicts noise
- Noise schedule:  $\{\beta_t\}$  and derived  $\{\alpha_t, \bar{\alpha}_t\}$
- Mathematical framework: Reverse process theory

What we need to figure out 

- How to start: Where does  $\mathbf{x}_T$  come from?
- How to step: What's the update rule for  $\mathbf{x}_t \rightarrow \mathbf{x}_{t-1}$ ?
- How many steps: Can we go faster than 1000 timesteps?
- How much randomness: Should sampling be deterministic or stochastic?

**Let's build the bridge from trained models to generated samples!**

Remember how we parameterized the reverse process?

The learned reverse transition

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_{\theta}(\mathbf{x}_t, t), \tilde{\sigma}_t^2 \mathbf{I})$$

Mean parameterization (from noise prediction)

$$\mu_{\theta}(\mathbf{x}_t, t) = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right)$$

Fixed variance

$$\tilde{\sigma}_t^2 = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$$

**Key insight:** We learn the mean, variance comes from the noise schedule! 

# DDPM Sampling: Following the Stochastic Path

The original sampling approach: embrace the randomness!

**Input:** Trained model  $\epsilon_\theta$ , timesteps  $T$

Sample  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

▷ Start from pure noise

**for**  $t = T, T - 1, \dots, 1$  **do**

    Predict noise:  $\hat{\epsilon} = \epsilon_\theta(\mathbf{x}_t, t)$

    Compute mean:  $\boldsymbol{\mu} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \alpha_t}} \hat{\epsilon} \right)$

**if**  $t > 1$  **then**

        Sample noise:  $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

$\mathbf{x}_{t-1} = \boldsymbol{\mu} + \sqrt{\tilde{\sigma}_t^2} \cdot \mathbf{z}$

**else**

$\mathbf{x}_0 = \boldsymbol{\mu}$

**end if**

▷ No noise in final step

**end for**


**Return:**  $\mathbf{x}_0$

**1000 steps later...** We have our generated sample! 

# Understanding Each Step of DDPM Sampling

1. Initialization: Start from chaos 

$\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  - Pure Gaussian noise

2. Noise prediction: What corruption was added? 

$\hat{\epsilon} = \epsilon_{\theta}(\mathbf{x}_t, t)$  - Use our trained network

3. Mean computation: Where should we go? 

Remove predicted noise to get expected previous state

4. Stochastic sampling: Add controlled randomness 

Don't just use the mean - add calibrated noise for diversity

5. Final step: Clean finish 

At  $t = 1$ , typically don't add noise (deterministic final step)

# The DDPM Trade-offs: Quality vs. Speed

## ✓ DDPM Strengths:

- **High quality:** Excellent samples with many steps
- **Diversity:** Stochastic sampling ensures variety
- **Theory:** Provably samples from correct distribution
- **Robustness:** Works across datasets and architectures

## ✗ DDPM Weaknesses:

- **Slow:** Requires 1000 function evaluations
- **Expensive:** Each step needs full network forward pass
- **Fixed:** Hard to trade quality for speed
- **Sequential:** Can't easily parallelize steps

**The bottleneck:** 1000 steps  $\times$  expensive network = too slow for real applications!

Can we do better?  Enter DDIM...



**What if we made sampling deterministic?**  
**The key realization... 💡**

What if we made sampling deterministic?

The key realization... 💡

DDIM's brilliant insight 🧠

DDPM defines ONE way to reverse the process...

But there are MANY reverse processes with the same marginals!

What if we made sampling deterministic?

The key realization... 💡

DDIM's brilliant insight 🧠

DDPM defines ONE way to reverse the process...

But there are MANY reverse processes with the same marginals!

What this means 🎯

- Same training, same network, same forward process
- But different (faster!) sampling procedure
- Maintains quality while dramatically improving speed
- **Can skip steps without retraining!**

This changes everything! 🚀

## The key mathematical insight behind DDIM...

### DDPM approach

Uses the Markovian reverse:  $q(\mathbf{x}_{t-1}|\mathbf{x}_t)$

### DDIM approach

Uses the more general:  $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$

*Condition on BOTH current state AND original data!*

### The constraint

We require the same marginals:  $q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$

**Result:** A family of reverse processes parameterized by  $\sigma_t!$  

# The Mathematical Surprise: $\sigma_t$ Emerges!

What happens when we solve the DDIM constraint problem?

The constraint we impose

Find  $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$  such that:

$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$  (same as DDPM)

The mathematical discovery 

When we solve this constraint (using Bayes rule + Gaussian algebra), we discover:

**The solution is not unique!**

There's a **family of valid solutions** parameterized by the variance  $\sigma_t^2$

# The Mathematical Surprise: $\sigma_t$ Emerges!

## The mathematical discovery

When we solve this constraint (using Bayes rule + Gaussian algebra), we discover:

**The solution is not unique!**

There's a **family of valid solutions** parameterized by the variance  $\sigma_t^2$

## DDIM's key insight

$$q_{\sigma}(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \mu(\mathbf{x}_t, \mathbf{x}_0), \sigma_t^2 \mathbf{I})$$

$\sigma_t^2$  can be *any* value in range  $[0, \beta_t]$  - it's a **free parameter**!

**Surprise:** This gives us unexpected control over randomness! 

# When $\sigma_t = 0$ : What Happens? ?

**The deterministic case - no randomness!**

## The big question

If we remove all randomness from the reverse process, how do we actually update from  $\mathbf{x}_t$  to  $\mathbf{x}_{t-1}$ ?

## DDIM's key insight 🎯

**Use the noise schedule structure!**

We know:  $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$

If we can estimate  $\mathbf{x}_0$  and  $\epsilon$ , we can construct  $\mathbf{x}_{t-1}$  directly!

**Next:** Let's see how this construction works step by step... →

Both DDPM and DDIM rely on this fundamental relationship:


## The universal noise schedule formula

At *any* timestep  $s$ , a noisy image has exactly this structure:

$$\mathbf{x}_s = \sqrt{\bar{\alpha}_s} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_s} \epsilon$$

## What this means

- $\mathbf{x}_0$ : The original clean image
- $\epsilon$ : The noise vector
- $\sqrt{\bar{\alpha}_s}$ : How much "signal" remains at timestep  $s$
- $\sqrt{1 - \bar{\alpha}_s}$ : How much "noise" is present at timestep  $s$

**DDIM's insight:** If we know  $\mathbf{x}_0$  and  $\epsilon$ , we can construct *any* timestep! 



# Step 1: Reverse Engineer the Components

Given  $\mathbf{x}_t$ , figure out what  $\mathbf{x}_0$  and  $\epsilon$  should be

## Estimate the noise

Use our trained neural network:

$$\hat{\epsilon} = \epsilon_{\theta}(\mathbf{x}_t, t)$$

## Solve for the clean image

Rearrange the noise schedule formula to find  $\mathbf{x}_0$ :

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \hat{\epsilon}$$

$\Downarrow$

$$\hat{\mathbf{x}}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \hat{\epsilon}}{\sqrt{\bar{\alpha}_t}}$$

Now we have both pieces:  $\hat{\mathbf{x}}_0$  and  $\hat{\epsilon}$  

## Step 2: Construct the Target Timestep

Use the same noise schedule formula to build  $\mathbf{x}_{t-1}$

Apply the noise schedule at timestep  $t - 1$

We want:  $\mathbf{x}_{t-1} = \sqrt{\bar{\alpha}_{t-1}}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_{t-1}}\epsilon$

Substitute our estimates

Replace  $\mathbf{x}_0$  and  $\epsilon$  with our predictions:

$$\mathbf{x}_{t-1} = \sqrt{\bar{\alpha}_{t-1}}\hat{\mathbf{x}}_0 + \sqrt{1 - \bar{\alpha}_{t-1}}\hat{\epsilon}$$

Why this works

- We're using the **same underlying clean image**  $\hat{\mathbf{x}}_0$
- We're using the **same noise vector**  $\hat{\epsilon}$
- We're just **adjusting the noise level** for timestep  $t - 1$

# The Complete DDIM Formula: Putting It Together


## Substitute Step 1 into Step 2:

The DDIM update rule ( $\eta = 0$ )

$$\mathbf{x}_{t-1} = \underbrace{\sqrt{\bar{\alpha}_{t-1}} \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_{\theta}(\mathbf{x}_t, t)}{\sqrt{\bar{\alpha}_t}}}_{\hat{\mathbf{x}}_0 \text{ (estimated clean image)}} + \sqrt{1 - \bar{\alpha}_{t-1}} \epsilon_{\theta}(\mathbf{x}_t, t)$$

The logic is crystal clear

- 1 **Estimate:** What are the clean image and noise?
- 2 **Construct:** Use these estimates to build the next timestep
- 3 **Consistency:** Same noise, different noise level

**Key insight:** We're not “removing then adding” noise - we're **reconstructing** the trajectory! 

# DDIM vs DDPM: Fundamentally Different Approaches

Two completely different philosophies:

## DDPM Philosophy ✂

### “Remove Noise Gradually”

- Take tiny steps
- Remove a bit of noise each time
- Hope randomness averages out
- Many careful iterations

*Like chiseling a sculpture*

## DDIM Philosophy 🎯

### “Reconstruct the Path”

- Estimate the clean image
- Use noise schedule structure
- Jump directly to target
- Fewer confident steps

*Like using a blueprint*

**Result:** DDIM achieves similar quality with 10-50x fewer steps! ★

# DDIM Sampling Algorithm: Speed and Elegance

**Input:** Trained model  $\epsilon_\theta$ , timestep subset  $\{\tau_1, \tau_2, \dots, \tau_S\}$

Sample  $\mathbf{x}_{\tau_S} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

**for**  $i = S, S - 1, \dots, 1$  **do**

$t = \tau_i, s = \tau_{i-1}$  (where  $\tau_0 = 0$ )

    Predict noise:  $\hat{\epsilon} = \epsilon_\theta(\mathbf{x}_t, t)$

    Predict  $\mathbf{x}_0$ :  $\hat{\mathbf{x}}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \hat{\epsilon}}{\sqrt{\bar{\alpha}_t}}$

    Update:  $\mathbf{x}_s = \sqrt{\bar{\alpha}_s} \hat{\mathbf{x}}_0 + \sqrt{1 - \bar{\alpha}_s} \hat{\epsilon}$

**end for**

**Return:**  $\mathbf{x}_0$

## Key differences from DDPM:

- **Timestep subset:** Use only  $S \ll T$  steps (e.g., 50 instead of 1000)
- **No random sampling:** Completely deterministic
- **Direct prediction:** Explicitly predict  $\mathbf{x}_0$  then re-noise

# DDPM vs. DDIM: The Ultimate Showdown

When should you use which method?

## ✓ Use DDPM when:

- Sample diversity is crucial
- Computational time isn't a constraint
- Need highest possible quality
- Theoretical guarantees matter
- Working with research/small-scale

## Typical settings:

- 1000 steps
- Stochastic sampling
- Research applications

## 🚀 Use DDIM when:

- Speed is important
- Need deterministic generation
- High-resolution images
- Real-time applications
- Production deployment

## Typical settings:

- 20-50 steps
- $\eta = 0$  or small
- Production applications

**Modern default:** DDIM with 50 steps for most applications! ★

# Preview: What's Coming Next

We can now generate unconditional samples...

But what if we want **CONTROL**? 

## Conditional generation challenges

- Generate specific classes: "Give me a cat image"
- Text-to-image: "A sunset over mountains"
- Style control: "In the style of Van Gogh"
- Spatial control: "Put the object here"





## Next lecture preview: Guided Sampling

- **Classifier guidance:** Use external classifiers to steer generation
- **Classifier-free guidance:** The breakthrough behind Stable Diffusion
- **Text conditioning:** How CLIP embeddings guide the process
- **Advanced control:** ControlNet and spatial conditioning

From random generation to controllable creation! 

# Summary: Sampling Mastery Achieved

## What we've learned today

-  **DDPM sampling:** Stochastic reverse process with theoretical guarantees
-  **DDIM sampling:** Deterministic alternative enabling dramatic acceleration
-  **Trade-offs:** Quality, speed, and diversity relationships
-  **Implementation:** Practical considerations for robust sampling

## Key practical insights





- Modern default: DDIM with 20-50 steps for most applications
- Use  $\eta$  parameter to control stochasticity vs. speed trade-off
- Deterministic sampling enables reproducibility and faster iteration
- Higher-order methods and optimizations provide further improvements


**We can now turn trained diffusion models into practical sample generators!** 



# Next Session Preview: Conditional Generation & Guidance

We'll explore how to control what diffusion models generate:

- How do we condition generation on classes, text, or images? 
- What is classifier guidance and why does it work? 
- How does classifier-free guidance eliminate the need for separate classifiers? 
- What makes text-to-image generation possible? 

**From unconditional noise-to-data generation  
to controllable, guided creation! **

**Ready to add steering wheels to our diffusion models? **