

Diffusion Models: Theory and Applications

Research Discussion: Generative Models for Information Retrieval, NLP, and RAG Systems

Shubham Chatterjee

Missouri University of Science and Technology, Department of Computer Science

June 13, 2025

Identifying High-Impact Research Opportunities

Session Goals

- Explore cutting-edge applications of generative models
- Identify novel research directions in IR/NLP/RAG
- Discuss fundable research proposals
- Connect theory to practical applications
- Foster cross-disciplinary collaborations

From established techniques to breakthrough innovations! 

The Convergence Opportunity

Why NOW is the perfect time for this research...

Mature Technologies

- Diffusion models (DDPM, DDIM, CFG)
- Variational Autoencoders (VAE, β -VAE)
- Generative Adversarial Networks
- Transformer architectures
- Large Language Models

Emerging Needs

- Intelligent information retrieval
- Contextual document generation
- Personalized content systems
- Knowledge-grounded generation
- Multimodal understanding

The intersection creates unprecedented research opportunities!

Beyond traditional keyword matching...

Core Innovation

Idea: Use generative models to synthesize retrieval results rather than just ranking existing documents.

Idea: Diffusion-Powered Query Understanding

- **Traditional Approach:**

- User query → keyword extraction → document matching
- Limited understanding of implicit intent
- Struggles with ambiguous or incomplete queries

- **New:**

- **Query Embedding Space:** Learn latent representations of search intents
- **Conditional Generation:** Generate query variations and clarifications
- **Intent Diffusion:** Sample from intent distribution given partial queries
- **Result Synthesis:** Generate comprehensive answers from intent understanding

Current RAG Limitations

- Static retrieval strategies
- Limited context integration
- Poor handling of contradictory sources
- Lack of uncertainty quantification

Generative Model Solutions

- **Diffusion RAG:** Generate retrieval queries iteratively
- **Quantum RAG:** Use ideas from quantum physics to develop advanced RAG systems

Idea 1: Diffusion RAG - Iterative Context Refinement

From one-shot retrieval to progressive context discovery...

Traditional RAG Limitations

- Fixed retrieval strategy regardless of query complexity
- Single-pass document selection without refinement
- Limited exploration of information space
- No uncertainty quantification in retrieval decisions

Diffusion RAG Innovation

Core Idea: Use diffusion processes to iteratively refine retrieval contexts, starting from broad noisy representations and progressively denoising to precise, relevant information.

Idea 1: Diffusion RAG - Iterative Context Refinement

Diffusion RAG Innovation

```
Input: Query  $q$ , Knowledge corpus  $\mathcal{K}$   
// Initialize with noisy context representation  
 $c_T \sim \mathcal{N}(0, I)$  // Random noise  
// Progressive denoising steps  
for  $t = T$  to 1 do  
     $c_{t-1} = \text{DiffusionStep}(c_t, q, \mathcal{K}, t)$   
     $\text{relevant\_docs}_t = \text{Sample}(c_{t-1}, \mathcal{K})$   
end for  
// Generate final response  
 $\text{response} = \text{LLM}(c_0, q)$ 
```

Key Advantages ★

- **Adaptive Exploration:** Discovers relevant context through iterative refinement
- **Uncertainty Aware:** Natural probabilistic estimates of relevance
- **Query Evolution:** Information need clarifies through the diffusion process

Idea 2: Quantum RAG Systems

Leveraging quantum computing for exponentially enhanced retrieval...

Quantum Computing Principles in RAG

- **Superposition:** Queries and documents exist in multiple states simultaneously
- **Entanglement:** Non-local correlations between related information pieces
- **Interference:** Constructive/destructive patterns naturally filter relevance
- **Amplitude Amplification:** Boost probability of finding highly relevant content

Quantum RAG Architecture

- **Quantum Query Encoding:** Map natural language to quantum state representations
- **Superposed Document Space:** All documents exist in quantum superposition
- **Entangled Context Networks:** Related concepts share quantum correlations
- **Quantum Search Algorithms:** Grover's algorithm variants for retrieval
- **Measurement-Based Selection:** Collapse to most relevant information

Potential Breakthroughs

- **Exponential Speedup:** $O(\sqrt{N})$ search complexity vs. classical $O(N)$
- **Hidden Correlations:** Discover non-obvious relationships through entanglement
- **Parallel Semantic Exploration:** Simultaneously explore multiple meanings
- **Natural Ambiguity Handling:** Superposition represents uncertain queries

Research Theme 3: Multimodal Knowledge Integration

The Multimodal Challenge

- Most knowledge exists across modalities (text, images, audio, video)
- Current systems struggle with cross-modal reasoning
- Limited integration of visual and textual information

Generative Solutions

- **Cross-Modal VAEs:** Shared latent spaces for text and images
- **Multimodal Diffusion:** Joint generation of text and visual content
- **GAN-Based Translation:** Convert between modalities seamlessly
- **Hierarchical Fusion:** Multi-level integration strategies

Applications

- Scientific document understanding with figures
- Visual question answering from documents
- Automatic illustration generation for text
- Cross-modal information retrieval

Research Theme 4: Personalized Knowledge Systems

The Personalization Opportunity

- Every user has unique knowledge backgrounds
- Information needs vary by expertise level
- Context depends on personal and professional goals

User Modeling with VAEs

- Learn latent user representations
- Capture reading preferences
- Model knowledge evolution
- Predict information needs

Content Adaptation with Diffusion

- Generate personalized explanations
- Adapt complexity levels
- Create custom examples
- Synthesize relevant analogies

Research Challenges

- How to model user knowledge accurately?
- What constitutes effective personalization?
- How to balance personalization with serendipitous discovery?

Research Problem

How can we create educational RAG systems that adapt to individual learning styles and knowledge levels?

Technical Approach

- **Student Modeling:** VAE for learning style representation
- **Content Generation:** Diffusion models for personalized explanations
- **Knowledge Tracking:** GAN-based assessment of understanding
- **Curriculum Planning:** Reinforcement learning for optimal sequencing

Research Theme 5: Factual Accuracy and Hallucination Control

The Hallucination Problem

- Generative models often produce plausible but false information
- Traditional evaluation metrics don't capture factual accuracy
- Users may not detect subtle inaccuracies

Novel Research Directions

- **Uncertainty-Aware Generation:** VAEs with explicit uncertainty modeling
- **Fact-Grounded Diffusion:** Constrained generation with knowledge bases
- **Adversarial Fact Checking:** GANs for detecting false information
- **Retrieval-Constrained Generation:** Hard constraints from reliable sources

Evaluation Innovations

- Automated fact-checking metrics
- Human-AI collaborative evaluation
- Uncertainty calibration measures
- Source attribution techniques

Research Theme 6: Efficient and Scalable Systems

Scalability Challenges

- Diffusion models require many inference steps
- VAE encoding/decoding adds latency
- GAN training can be unstable at scale
- Real-time requirements for interactive systems

Efficiency Innovations

- Few-step diffusion sampling
- Quantized VAE representations
- Distilled generative models
- Hierarchical caching strategies

System Optimizations

- Edge deployment techniques
- Distributed inference
- Adaptive model selection
- Progressive generation

Research Opportunities

- Novel sampling algorithms for faster generation
- Compression techniques for generative models
- Hardware-software co-design for inference

Research Theme 7: AI for Indian Communities

Unique Opportunities in Indian Context

- Large, diverse population with varied literacy levels
- Multilingual environment (22 official languages + dialects)
- Resource constraints requiring efficient solutions
- Strong mobile-first technology adoption
- Rich cultural and traditional knowledge systems

Research Focus Areas

- **Agricultural Intelligence:** Supporting 600+ million farmers
- **Educational Technology:** Enhancing learning for 250+ million students
- **Healthcare AI:** Improving access in resource-limited settings
- **Multilingual Systems:** Bridging language barriers
- **Cultural Preservation:** Digitizing traditional knowledge

Massive scale + Real impact = World-class research! 

Application 1: AI-Powered Agricultural Assistance

The Challenge

- 86% of farmers are small and marginal (< 2 hectares)
- Limited access to agricultural experts and extension services
- Climate change creating unpredictable growing conditions
- Language barriers in accessing technical information

Generative AI Solutions

- **Crop Advisory VAE:** Generate personalized farming recommendations
- **Weather-Crop Diffusion:** Predict optimal planting strategies
- **Multilingual RAG:** Answer farming questions in local languages
- **Visual Crop Diagnosis:** Generate explanations from crop images

Research Innovation

Multimodal Agricultural Advisor: Combine satellite imagery, weather data, soil sensors, and farmer queries to generate contextual farming advice in regional languages.

Application 2: Intelligent Teaching Assistant for Indian Schools

Educational Challenges in India

- High student-teacher ratios (30:1 to 50:1 in many schools)
- Multi-grade classrooms with mixed ability levels
- Teachers need support in creating engaging content
- Language transition from mother tongue to English

AI-Powered Solutions

- **Adaptive Content Generation:** Create lessons for different learning levels
- **Multilingual Explanation VAE:** Generate concepts in multiple languages
- **Question Generation Diffusion:** Create practice problems and assessments
- **Cultural Context RAG:** Incorporate local examples and stories

Research Innovation

Culturally-Aware Educational AI: Generate teaching materials that incorporate local cultural references, examples from student's environment, and traditional knowledge systems.

Application 3: Rural Healthcare AI Assistant

Healthcare Access Challenge

- Doctor-patient ratio: 1:1456 (WHO recommends 1:1000)
- 65% of population lives in rural areas with limited healthcare
- ASHA workers need decision support tools
- Language barriers in medical communication

Generative AI for Health

- **Symptom-to-Advice VAE:** Generate preliminary health guidance
- **Medical Image Diffusion:** Enhance low-quality diagnostic images
- **Treatment Plan RAG:** Contextualize medical protocols for local settings
- **Health Education Generation:** Create prevention awareness content

Ethical Considerations

- Clear limitations and referral protocols
- Integration with existing healthcare systems
- Privacy protection for health data
- Validation with medical professionals

Beyond single models: Intelligent agent collaboration...

The Agent Revolution

- Moving from passive models to proactive agents
- Agents that can plan, execute, and adapt
- Multi-agent collaboration for complex tasks
- Integration of reasoning, action, and learning

Research Opportunities

- **Agent Architecture Design:** Novel frameworks for autonomous behavior
- **Multi-Agent Coordination:** Emergent intelligence from agent interactions
- **Agent Training Methods:** Beyond supervised learning approaches
- **Human-Agent Collaboration:** Seamless integration with human workflows

Technical Challenges

- Long-term planning and goal decomposition
- Safe exploration in real-world environments
- Communication and coordination protocols
- Scalability of multi-agent systems

Novel Agent Training Paradigms

Traditional Limitations

- Supervised learning requires extensive labeled data
- Reinforcement learning sample inefficiency
- Limited transfer across domains and tasks
- Brittleness in novel situations

Innovative Training Methods

- **Synthetic Experience Generation:** Use diffusion models to create training scenarios
- **Curriculum via VAE:** Progressive skill development through latent space exploration
- **Adversarial Agent Training:** GAN-like competition between agent populations
- **Constitutional AI for Agents:** Value-aligned agent behavior through self-reflection

Breakthrough Approach

Generative Agent Training: Use diffusion models to generate infinite diverse training environments and scenarios, enabling agents to learn robust behaviors.

Idea 1: Multi-Agent Agricultural Extension System

Agent Ecosystem Design

- **Weather Agent:** Monitors and predicts weather patterns
- **Crop Agent:** Tracks crop health and growth stages
- **Market Agent:** Analyzes prices and demand trends
- **Advisory Agent:** Integrates insights for farmer recommendations
- **Communication Agent:** Delivers advice in appropriate language/format

Agent Collaboration

- Agents share observations and predictions
- Distributed decision-making for optimal farm management
- Emergent strategies from agent interactions
- Adaptive coordination based on changing conditions

Training Innovation

Train agents in simulated farming environments generated by diffusion models, covering diverse crops, weather patterns, and economic conditions across India.

Idea 2: Curriculum Learning for Educational Agents

Adaptive Teaching Challenge

- Students have different learning paces and styles
- Need to balance challenge and accessibility
- Must maintain engagement while ensuring comprehension
- Require personalization at scale

VAE-Driven Curriculum Design

- **Student State Encoding:** Map knowledge and skill levels to latent space
- **Optimal Path Generation:** Sample learning trajectories from latent space
- **Content Adaptation:** Generate materials matching student state
- **Progress Tracking:** Update student representation based on performance

Research Innovation

Generative Curriculum AI: Use diffusion models to generate personalized learning sequences that adapt in real-time to student progress and preferences.

Idea 3: Constitutional AI for Value-Aligned Agents

Cultural Alignment Challenge

- AI systems must respect diverse Indian cultural values
- Balance individual needs with community welfare
- Incorporate concepts like *ahimsa*, *seva*, and *dharma*
- Handle multi-religious and multi-cultural contexts

Constitutional Training Process

- **Value Encoding:** Represent cultural principles in training objectives
- **Self-Reflection:** Agents evaluate their actions against value systems
- **Community Feedback:** Incorporate local community input in training
- **Continuous Alignment:** Ongoing refinement based on cultural feedback

Research Questions

- How to formalize cultural values for AI training?
- Can agents learn appropriate behavior from community interactions?
- How to balance conflicting values in decision-making?

Idea 4: Self-Improving Agent Architectures

Self-Improvement Paradigm

- Agents analyze their own performance and identify weaknesses
- Generate additional training data based on failure cases
- Modify their own architectures and learning strategies
- Engage in meta-learning and strategy optimization

Implementation Approaches

- **VAE-Based Skill Discovery:** Learn new capabilities in latent space
- **Diffusion for Experience Generation:** Create challenging training scenarios
- **GAN-Style Skill Competition:** Agents compete to develop better strategies
- **Neural Architecture Search:** Evolve better agent architectures

Research Theme 10: LLM-Based Content Generation and Curation

Content Creation Challenges

- Maintaining factual accuracy in generated content
- Balancing creativity with reliability
- Incorporating real-time information and updates
- Personalizing content for different audiences

LLM-Powered Solutions

- **Fact-Checked Generation:** Real-time verification against knowledge bases
- **Style-Adaptive Writing:** Content tailored to audience and purpose
- **Collaborative Authoring:** Human-AI co-creation workflows
- **Multi-Source Synthesis:** Combining diverse information sources
- **Version Control for Ideas:** Tracking content evolution and sources

Applications

- Automated technical documentation generation
- Personalized educational material creation
- News article synthesis from multiple sources
- Research paper summarization and analysis

Research Theme 11: Efficient LLM Architectures and Deployment ⚡

Efficiency Challenges 📊

- Large model sizes require significant computational resources
- High latency limits real-time applications
- Memory constraints on edge devices
- Energy consumption and environmental impact

Model Compression 🔄

- Knowledge distillation techniques
- Pruning and quantization methods
- Low-rank approximation
- Dynamic model sizing

Inference Optimization ⚡

- Speculative decoding
- Batched inference strategies
- Caching mechanisms
- Edge deployment techniques

Novel Research Directions 🎯

- **Mixture of Experts (MoE):** Activate only relevant model parts
- **Progressive Generation:** Start with fast approximations, refine as needed
- **Adaptive Precision:** Dynamic bit-width based on importance
- **Federated Inference:** Distribute computation across devices

Innovation: Hierarchical LLM Architecture for Efficiency

The Efficiency Trade-off

How can we maintain LLM quality while dramatically reducing computational requirements?

Proposed Hierarchical System

- **Router LLM (Small):** Classifies query complexity and routes appropriately
- **Fast LLM (Medium):** Handles simple queries with low latency
- **Expert LLM (Large):** Processes complex queries requiring deep reasoning
- **Specialist LLMs:** Domain-specific models for specialized tasks
- **Confidence Monitor:** Determines when to escalate to larger models

Key Innovations

- Dynamic routing based on query analysis and user context
- Progressive refinement: start fast, improve if needed
- Learned routing strategies that improve over time
- Uncertainty-based escalation protocols

Goal: 80% performance at 20% computational cost!

Research Theme 12: LLMs for Code Generation and Software Engineering

Transforming software development with intelligent code generation...

Code Generation Challenges ⚠️

- Generating syntactically correct and semantically meaningful code
- Understanding complex requirements and specifications
- Maintaining code quality, security, and performance
- Handling diverse programming languages and frameworks

LLM-Powered Development 🧠

- **Requirement-to-Code Translation:** Natural language to implementation
- **Code Completion and Suggestion:** Context-aware programming assistance
- **Bug Detection and Fixing:** Automated debugging and repair
- **Code Review and Documentation:** Intelligent analysis and explanation
- **Test Generation:** Automated test case creation and validation

Research Opportunities 💡

- Multi-language code translation and optimization
- Integration with software development workflows
- Code security analysis and vulnerability detection

Research Theme 13: Continual Learning and Knowledge Updates

Knowledge Obsolescence Problem

- LLMs trained on static datasets become outdated
- New information emerges constantly in dynamic domains
- Catastrophic forgetting when learning new knowledge
- Expensive retraining for knowledge updates

Continual Learning Solutions

- **Elastic Weight Consolidation:** Protect important model parameters
- **Memory-Augmented Networks:** External memory for new knowledge
- **Progressive Neural Networks:** Add capacity for new tasks
- **Knowledge Distillation Updates:** Efficient transfer of new information
- **Retrieval-Based Updates:** External knowledge without retraining

Research Challenges

- How to detect when knowledge needs updating?
- What constitutes reliable new information sources?
- How to handle conflicting information gracefully?
- Can we develop self-updating LLM architectures?

Research Theme 14: Explainable and Interpretable LLMs

Understanding what LLMs know and how they reason...

Interpretability Challenges

- LLMs are black boxes with billions of parameters
- Difficult to understand decision-making processes
- Limited ability to verify knowledge sources
- Challenges in debugging model behavior

Explainability Approaches

- **Attention Visualization:** Understanding focus mechanisms
- **Probing Studies:** What knowledge is encoded where?
- **Activation Patching:** Causal intervention analysis
- **Natural Language Explanations:** LLMs explaining their reasoning
- **Concept Bottleneck Models:** Interpretable intermediate representations

Applications

- Medical diagnosis explanation and validation
- Legal reasoning transparency and accountability
- Educational feedback and tutoring systems
- Scientific hypothesis generation and verification

Innovation: Self-Explaining LLM Architecture

The Explanation Challenge

How can we build LLMs that naturally provide explanations for their reasoning without sacrificing performance?

Proposed Architecture

- **Reasoning Module:** Explicit step-by-step thinking process
- **Evidence Tracker:** Identifies and cites relevant information sources
- **Confidence Estimator:** Quantifies certainty in different parts of response
- **Explanation Generator:** Creates natural language explanations of reasoning
- **Verification Module:** Cross-checks reasoning for consistency

Training Innovations

- Multi-task learning with explanation generation
- Reinforcement learning from human explanation feedback
- Self-supervised learning from reasoning traces
- Constitutional training for explanation quality

Goal: Trustworthy AI through transparent reasoning!

Research Project: Domain-Adaptive LLM Framework

Project Vision

Develop a framework that enables rapid adaptation of general-purpose LLMs to specialized domains (medical, legal, scientific) with minimal training data and computational resources.

Technical Approach

- **Domain Knowledge Extraction:** Automatically identify key concepts and relationships
- **Few-Shot Domain Adaptation:** Efficient fine-tuning with limited examples
- **Knowledge Graph Integration:** Incorporate structured domain knowledge
- **Transfer Learning Strategies:** Leverage knowledge from related domains
- **Performance Monitoring:** Continuous evaluation and improvement

Applications

- Legal document analysis and contract review
- Medical literature summarization and diagnosis support
- Scientific paper understanding and hypothesis generation
- Financial analysis and risk assessment
- Engineering design and optimization

Impact: Democratize specialized AI across industries!

Challenge

Training high-quality LLMs for Indian regional languages requires distributed data collection while preserving privacy and cultural sensitivity.

Federated Learning Approach

- **Decentralized Training:** Train on local devices without data sharing
- **Language-Specific Adapters:** Efficient multilingual model architectures
- **Cultural Knowledge Preservation:** Maintain regional nuances and context
- **Privacy-Preserving Aggregation:** Secure model update mechanisms
- **Resource-Efficient Methods:** Optimize for mobile and edge devices

Research Innovations

- Novel aggregation algorithms for heterogeneous language data
- Cross-lingual transfer learning with privacy constraints
- Cultural bias detection and mitigation techniques
- Adaptive model architectures for resource-constrained environments

Impact: Bridge digital divide through inclusive AI!

Cross-Disciplinary Research Opportunities

Collaborations beyond computer science...

Digital Humanities

- Historical document generation
- Literary style transfer
- Cultural bias analysis
- Archival knowledge extraction

Psychology

- Cognitive load modeling
- Attention mechanism studies
- Learning preference analysis
- Memory and recall optimization

Education

- Adaptive learning systems
- Automated content generation
- Student modeling
- Assessment innovation

Library Science

- Knowledge organization
- Information architecture
- User experience design
- Collection development

Interdisciplinary research leads to breakthrough innovations!

Sample Research Timeline

Year 1: Foundation Building

- Literature review and gap analysis
- Preliminary experiments and proof-of-concept
- Baseline implementations and datasets
- First conference submissions (workshops)

Year 2: Core Innovation

- Novel algorithm development
- Comprehensive experimental evaluation
- Comparison with state-of-the-art
- Major conference submissions (NeurIPS, ICML, ACL)

Year 3: Impact and Expansion

- Real-world deployment and user studies
- Open-source release and community adoption
- Journal publications and survey papers
- Grant proposals for follow-up research

Conference Venues

- **ML:** NeurIPS, ICML, ICLR
- **NLP:** ACL, EMNLP, NAACL
- **IR:** SIGIR, WWW, CIKM
- **AI:** AAAI, IJCAI
- **Systems:** OSDI, SOSP (for scalability work)

Journal Options

- **High Impact:** Nature Machine Intelligence, Science Robotics
- **ML Journals:** JMLR, Machine Learning
- **NLP Journals:** TACL, Computational Linguistics
- **IR Journals:** TOIS, Information Retrieval

Publication Pipeline →

- Workshop papers for early feedback
- Conference papers for peer review and visibility
- Journal papers for comprehensive contributions
- Survey papers to establish thought leadership

Implementation Roadmap

Immediate Actions (Next 3 Months) ▶

- Form research team with complementary expertise
- Conduct comprehensive literature review
- Set up experimental infrastructure
- Apply for seed funding or equipment grants

Medium-term Goals (6-12 Months) ◎

- Complete baseline implementations
- Submit to relevant workshops
- Establish industry collaborations
- Recruit graduate students and postdocs

Long-term Vision (1-3 Years) ◉

- Build recognized research program
- Secure major funding (NSF CAREER, etc.)
- Establish international collaborations
- Create lasting impact in the field

Building a research network...

Academic Partnerships

- Joint student supervision
- Shared computational resources
- Complementary expertise exchange
- International research visits

Industry Connections

- Internship programs
- Real-world problem validation
- Access to large-scale datasets
- Technology transfer opportunities

Conference Networking

- Workshop organization
- Special session proposals
- Panel discussions
- Social media engagement

Open Source Community

- Code repositories and libraries
- Benchmark datasets
- Reproducible research
- Community challenges

Research is a team sport - build your network! 

Engaging students in cutting-edge research...

Undergraduate Opportunities

- Independent study projects
- Summer research experiences (REU)
- Honors thesis programs
- Competition teams (Kaggle, DrivenData)

Graduate Student Projects

- **MS Projects:** Implementation and evaluation studies
- **PhD Research:** Novel algorithmic contributions
- **Collaborative Projects:** Industry partnerships
- **Cross-disciplinary Work:** Domain-specific applications

Mentorship Best Practices

- Start with manageable, well-defined projects
- Provide regular feedback and guidance
- Encourage conference participation
- Foster collaborative research culture

From research to real-world impact...

Commercialization Pathways

- Licensing to existing companies
- Startup company formation
- Open-source community adoption
- Government and non-profit partnerships

Considerations

- Intellectual property protection
- University policies and agreements
- Funding source restrictions
- Ethical implications of commercialization

Future Directions and Emerging Trends

Where is the field heading?

Technical Trends

- Multimodal foundation models
- Efficient training and inference
- Continual and few-shot learning
- Neuro-symbolic integration
- Quantum-enhanced algorithms

Application Trends

- Scientific discovery acceleration
- Personalized education at scale
- Creative industry transformation
- Healthcare and medical AI
- Environmental monitoring

Research Preparation

- Stay current with ArXiv and conference proceedings
- Participate in workshops and tutorials
- Engage with industry research labs
- Build flexible, modular research infrastructure

Position yourself at the forefront of emerging trends! 

Faculty and Student Research Planning

Discussion Topics

- Which research directions align with your interests?
- What computational resources do you need?
- How can we form collaborative research groups?
- What are the biggest technical challenges you foresee?

Action Items

- Identify potential research team members
- Define preliminary research questions
- Outline resource requirements
- Plan next steps and timelines

Turning Discussion into Action: Let's Draft Some Ideas!

Workshop Objective

- To collaboratively develop initial research proposals based on the themes discussed.
- Foster interaction and teamwork between faculty and students.
- Practice condensing complex ideas into a concise proposal format.
- Generate tangible starting points for future research projects.

Expected Outcome

A set of short (approx. 4-page) research proposals, co-authored by mixed groups, ready for discussion and refinement.

An opportunity to plant the seeds for innovative research!

Mini-Proposal Guidelines (4-Page Target)

Core Sections for the 4-Page Proposal

● Page 1: Title, Team, and Abstract

- Catchy Title & Names of Faculty/Student Collaborators.
- Abstract (approx. 250 words): Problem, core idea, approach, expected impact.

● Page 2: Problem Statement & Significance

- Clearly define the research problem.
- Explain its importance and relevance to frontier research areas (IR, NLP, RAG, GenAI).

● Page 3: Proposed Approach & Methodology

- Outline your novel generative model-based solution.
- Briefly describe key methods, techniques, or datasets you plan to use.

● Page 4: Expected Outcomes & Discussion Plan

- What are the anticipated results or contributions?
- Brief plan for evaluation or proof-of-concept.
- Points for in-class discussion.

Workshop Workflow: From Idea to Discussion

Phase 1: Ideation & Team Formation

- Form small groups (e.g., 3-4 students + 1-2 faculty member).
- Align based on shared interests from earlier discussions.
- Brainstorm and select a specific research question/idea.
- *Time: Approx. 30 minutes*

Phase 2: Collaborative Drafting

- Work together to draft the 4-page proposal using the guidelines.
- Assign sections or write collaboratively.
- Focus on getting the core ideas down.
- *Time: Approx. 90 minutes (or as an offline take-home task before next session)*

Phase 3: In-Class Proposal Discussion

- Each group briefly presents their proposal (e.g., 5-7 minutes).
- Open floor for constructive feedback, questions, and suggestions from all participants.
- Identify synergies, potential overlaps, or new collaborative avenues.
- *Time: Approx. 10-15 minutes per group in a subsequent session.*

Let's engage, create, and critique constructively!

Key Research Themes ✓

- 🔍 **Generative Information Retrieval:** Beyond traditional search
- **Neural Document Generation:** Creating content on-demand
- 🏗️ **Advanced RAG Architectures:** Smarter context integration
- 📄 **Multimodal Knowledge Integration:** Cross-modal understanding
- 👤 **Personalized Knowledge Systems:** Tailored information delivery
- ✅ **Factual Accuracy Control:** Reliable generative systems
- ⚡ **Efficient and Scalable Systems:** Practical implementations